

R and Python and 5*A & Nevai Search Engine Doc_Project

Ahmed Abobakr

Project : Student Management & AI-Powered Search Platform (5*A & Nevai)

1. Project Team

Order ID	Number	Name	Role in Project	Responsibilities
1	931230015	Ahmed Abobakr Hanafi	Project Manager / System Architect	Overall project supervision, architectural decisions, code review, team coordination
2	931230187	Ali Ez Alyan	Backend Lead Developer	API development, database management, security implementation, performance optimization
3	931230030	Ahmed Sabry Hamza	Frontend Developer	React/Vue UI development, user experience design, animations, dark/light theme implementation
4	931230037	Ahmed AbdulNasser Sayed	API Integration Specialist	Integrating 20+ search sources, data processing, error handling, response optimization
5	931230046	Ahmed Mohamed Mohamed	QA Engineer & Documentation	Testing, code review, project documentation, quality assurance

2. Project Goals and Technical Objectives

This project suite was engineered with distinct, measurable objectives to overcome common limitations found in monolithic educational software and standard search utilities. Each objective represents a targeted architectural decision aimed at enhancing system reliability, security, performance, and data accessibility.

Objective 1: High-Performance Fuzzy Search at Scale

- Problem Solved:* Inefficient student record retrieval across large datasets.
- Method Used:* PostgreSQL `pg_trgm` extension for trigram-based fuzzy matching.
- Architectural Outcome:* Sub-10ms fuzzy search capabilities across millions of characters.

Objective 2: Secure Session Management and Threat Mitigation

- Problem Solved:* Vulnerability to Cross-Site Scripting (XSS), Cross-Site Request Forgery (CSRF), and unauthorized access to sensitive student records.

- *Method Used:* Implementation of a robust JWT authentication strategy incorporating short-lived access tokens, long-lived refresh tokens, and strict transmission via `HttpOnly`, `Secure`, and `SameSite` cookies.
- *Architectural Outcome:* A stateless, secure session management layer neutralizing XSS vectors for session hijacking, while strict CORS and SameSite policies mitigate CSRF threats. Strict RBAC ensures role-delineated access.

Objective 3: High-Performance Data Retrieval Under Concurrent Load

- *Problem Solved:* Unacceptable response latencies during frequent, concurrent data retrieval requests and stale data reads.
- *Method Used:* Deployment of a Cache-Aside caching pattern utilizing an in-memory Redis instance with strict cache invalidation hooks on mutating operations.
- *Architectural Outcome:* Drastically reduced database read-load, achieving a target p95 latency of ~45ms for read-heavy endpoints with 100% data consistency.

Objective 4: Absolute Payload Structural Integrity

- *Problem Solved:* Inconsistent, malformed, or malicious data payloads leading to database schema violations or brittle API contracts.
- *Method Used:* Middleware-level request parsing and strict type validation using Pydantic schemas within the FastAPI ecosystem.
- *Architectural Outcome:* Guaranteed type-safe data ingestion. Malformed payloads are intercepted at the API boundary, returning `HTTP 422 Unprocessable Entity`.

Objective 5: Comprehensive System Observability

- *Problem Solved:* Lack of real-time system observability for detecting performance regressions or bottlenecks.
- *Method Used:* Integration of a Prometheus telemetry pipeline scraped via `/metrics`, visualized through Grafana dashboards, and structured JSON logging via Loguru.
- *Architectural Outcome:* Real-time tracking of endpoint latency, error rates, cache hit ratios, and audit trails.

Objective 6: Infrastructure Orchestration and Reproducibility

- *Problem Solved:* Complex, fragile local setup procedures and deployment topologies.
- *Method Used:* Multi-stage Docker containerization and unified orchestration (`run.sh`) via Docker Compose.
- *Architectural Outcome:* Guaranteed deterministic builds and zero-config local startup.

Objective 7: Automated Regression Defense

- *Problem Solved:* Unreliable code deployments where new features inadvertently break existing functionality.
- *Method Used:* Comprehensive automated test suite using `pytest`, featuring isolated fixtures and a dedicated test database.
- *Architectural Outcome:* Achieved 84% test coverage with 37 distinct automated test cases covering auth, RBAC, and CRUD.

Objective 8: Privacy-Focused AI Search Aggregation

- *Problem Solved:* Fragmented search experiences and privacy-invasive AI integrations requiring external API keys.
 - *Method Used:* Development of Nevai, an AI-powered multi-source search engine aggregating 20+ APIs, utilizing local LLM execution via WebAssembly (`web-llm`).
 - *Architectural Outcome:* A minimalist, privacy-first search engine capable of summarizing, comparing, and extracting data without sending user data to external servers.
-

3. Abstract

This project presents the architecture, implementation, and evaluation of a comprehensive, production-ready platform comprising three interconnected systems: a **Student Management System (SMS)**, the **Nevai AI-Powered Search Engine**, and the **5*A Task Management System**. Engineered to resolve the critical trilemma of scalability, security, and data integrity, the system employs a robust, distributed technology stack. The backend leverages the asynchronous capabilities of FastAPI (Python) backed by a normalized PostgreSQL database and an in-memory Redis caching layer. The client interface is a highly responsive React 19 Single Page Application orchestrated by Vite.

Key architectural innovations include a stringent, `HttpOnly` JWT-based authentication flow, Redis-backed cache-aside strategies, comprehensive audit logging with JSON diffing, and absolute boundary validation using Pydantic. Furthermore, the integration of `web-llm` enables privacy-preserving, local AI execution for advanced search summarization across 20+ aggregated data sources. The resulting deliverable is a highly available, fully containerized environment that guarantees a p95 latency of 45ms, maintains 84% automated test coverage, and provides a secure, role-delineated platform adaptable to varying educational and research environments.

4. Introduction

Context and Background

Modern educational institutions and knowledge workers operate dynamic, data-intensive environments. Administrators require real-time overviews of institutional health, instructors need reliable access to grading systems, students expect instantaneous access to academic records, and all users benefit from powerful, privacy-respecting search tools. Traditional systems often struggle to handle concurrent access peaks without compromising data security, experiencing performance degradation, or siloing knowledge retrieval.

System Evolution

To address these challenges, the project adopted a modern, distributed architectural approach. By decoupling the client application from the server logic and containerizing the infrastructure, the system eliminates traditional bottlenecks. The transition to a FastAPI backend allows for high-throughput, asynchronous request handling, while the integration of Redis ensures that read-heavy operations do not burden PostgreSQL. Concurrently, the Nevai search module aggregates over 20 external APIs, processed locally via WebAssembly-based AI to provide instant, private insights.

Core Contributions

- **Decoupled Client-Server Paradigm:** Enforcing a strict API contract between the React frontend and Python backend.

- **Advanced Authentication Flow:** Hardening session management by abstracting tokens to `HttpOnly` cookies with automated rotation.
- **Performant Data Access:** Implementing a Cache-Aside pattern with aggressive invalidation upon state mutation.
- **Privacy-First AI Search:** Running LLMs entirely in the browser to aggregate and summarize 20+ search sources without external data transmission.
- **Comprehensive Observability:** Embedding a Prometheus/Grafana telemetry pipeline for real-time micro-level performance metrics.

5. Methodology & Architecture

The architecture is divided into distinct functional modules to ensure separation of concerns.

Environment & Reproducibility

- **Tools/Libraries:** Docker, Docker Compose, Uvicorn (ASGI server), Vite, Bash (`run.sh`).
- **Architecture:** Multi-container orchestration isolating `db` (PostgreSQL), `redis`, `backend` (FastAPI), `frontend` (Nginx/React), `prometheus`, and `grafana`. Volume mapping ensures persistence.
- **Configuration:** Controlled via `.env` injection and explicit port mapping.
- **Module References:** `infra/docker-compose.yml`, `run.sh`, `fastapi_backend/Dockerfile`, `5A_Search/vite.config.js`.

Data Ingestion & Validation

- **Tools/Libraries:** FastAPI, Pydantic v2, Python typing, TypeScript interfaces.
- **Architecture:** Middleware-level interception. JSON payloads are cast against strict Pydantic models enforcing types, lengths, regex, and mandatory fields.
- **Validation Strategy:** Early rejection. Malformed payloads return `HTTP 422 Unprocessable Entity`.
- **Module References:** `fastapi_backend/app/schemas/`, `5A_Search/src/lib/api.ts`.

Core Processing Engine

- **Tools/Libraries:** FastAPI, Python 3.11, AsyncPG, SQLAlchemy, Alembic.
- **Architecture:** Controller/Service/Repository pattern. AsyncPG for scalable, non-blocking I/O.
- **Configuration:** RBAC enforced via FastAPI dependency injection (`Depends / OAuth2PasswordBearer`).
- **Module References:** `fastapi_backend/app/api/`, `fastapi_backend/app/services/`, `fastapi_backend/app/core/security.py`.

Validation & Metrics

- **Tools/Libraries:** Pytest, HTTPX, Prometheus, Grafana, Loguru.
- **Architecture:** Isolated test database via `pytest`. Custom middleware records latency/counters for Prometheus. Loguru handles structured JSON logging.
- **Metrics Table:**

Metric	Target Threshold	Current Status
API Latency (p95)	< 200ms	45ms
Cache Hit Ratio	> 80%	*Placeholder†

Metric	Target Threshold	Current Status
Test Coverage	> 80%	84% (37 Test Cases)

(† Requires production load generation to populate accurately)

Dashboard & Visualization


- **Tools/Libraries:** React 19, Vite, Zustand, TailwindCSS v4, Lucide React, React Router v6.
- **Architecture:** SPA with client-side routing. Zustand for lightweight global/persisted state. Tailwind CSS for utility-first styling. Includes productivity modules (Kanban, Notes, Calendar) and the Nevai Search Engine.
- **Module References:** 5A_Search/src/App.tsx , 5A_Search/src/store/ , 5A_Search/src/pages/ .

6. System Architecture & Tech Stack

Tech Stack

Layer	Technology
Frontend	React 19 + TypeScript + Vite + TailwindCSS v4 + Zustand + web-llm
Backend	FastAPI + Python 3.11 (Asynchronous)
Database	PostgreSQL 16 + SQLAlchemy (AsyncPG)
Cache	Redis 7 (Cache-Aside Pattern)
Logging	Loguru (Structured JSON) + Audit Logging
Metrics	Prometheus + Grafana + Custom API Monitoring
Security	JWT (HS256) + bcrypt + Password Strength Enforcement + XSS/CSRF Mitigation
AI / Search	web-llm (WebAssembly) + 20+ Aggregated APIs
DevOps	Docker Compose

Architecture Diagram



7. Module Deep Dive: Nevai AI Search Engine (5A_Search)

Overview

Nevai is an AI-powered, multi-source search engine built as a Single Page Application. It follows a minimalist, privacy-focused design philosophy that requires no external API keys from the end-user.

API Sources (20+ Aggregated Endpoints)

Category	Sources
General Knowledge	DuckDuckGo (Primary search, instant answers), SearXNG (Privacy meta-search), Wikipedia, Wikidata, DBpedia
Specialized Wikimedia	Wikiquote, Wikibooks, Wikivoyage, Wikisource, Wikiversity, Wikimedia Commons
Academic & Technical	Open Library (Books), arXiv (Scientific papers), Hacker News, GitHub, Dev.to
Social & Media	Reddit, iTunes, Deezer
Data & Services	CoinGecko (Crypto), OpenMeteo (Weather), RestCountries (Country data)

Search Categories

Type	Description
Web Search	General queries across multiple sources
News Search	Latest articles from news aggregators
Image Search	Visual content results
Instant Answers	Quick factual responses
Definitions	Word meanings and explanations
Weather	Location-based weather data
Time/Date	World clock and timezone data
Translations	Multi-language translation support

AI-Powered Features (Local LLM)

- **WebAssembly Execution:** Runs entirely in the browser using `web-llm`. No data is sent to external AI servers.
- **Privacy-Preserving:** Context-aware AI chat that keeps user queries strictly local.
- **Capabilities:** Generates search result summaries, answers follow-up questions, explains complex topics, compares search results, and extracts key information.

Security Implementation

Area	Implementation Details
Input Sanitization	XSS prevention, SQL injection protection, script tag filtering, URL validation.

Area	Implementation Details
Console Security	Console locking mechanism in production, warning messages for self-XSS, development mode detection.
Security Headers	Content Security Policy (CSP) checking, HTTPS enforcement, Referrer policy validation.

User Experience & Performance

- Features:** URL-bookmarkable searches (query parameters), browser history integration, local storage for preferences, dark mode with system preference detection, mobile-responsive design, PWA-ready.
- Performance:** Code splitting, lazy loading of components, optimized bundle size, fast initial load time.

8. Results & Discussion

Interpretation of Results:

The architectural shift to a decoupled FastAPI and React stack has yielded significant improvements. The Redis caching integration successfully offloaded repetitive read queries from PostgreSQL, resulting in a highly stable p95 latency of 45ms. Security posture is vastly improved through `HttpOnly` JWTs and strict Pydantic schema validation. Furthermore, the client-side integration of `web-llm` successfully demonstrated that complex AI summarization can occur locally, eliminating API latency and privacy concerns associated with external LLM providers.

Architectural Limitations:

The architecture relies on a single-instance PostgreSQL database and a standalone Redis cache, which presents potential single points of failure under enterprise-level loads (mitigated by Redis Sentinel/Cluster in future iterations). Additionally, the local AI model (`web-llm`) requires an initial download and depends on client hardware capabilities (WebAssembly support, available RAM), which may result in varied performance on older devices.

Strategic Implications and Future Scope:

The decoupled, containerized nature of the system allows the Nevai search module to be extracted into a standalone browser extension or integrated into third-party institutional portals. The distinct API boundaries facilitate straightforward extraction of domains into microservices. Deploying a native mobile application via React Native would require zero modifications to the core processing engine.

9. Reproducibility & Quick Start

To guarantee precise reproducibility, the system relies exclusively on automated shell scripts and Docker orchestration.

Exact Rebuild Steps

- Environment Preparation:** Ensure Docker Engine (v20+) and Docker Compose (v2+) are installed. Ensure ports 80, 8000, 5432, 6379, 9090, and 3000 are available.
- Repository Initialization:** Clone the repository and navigate to the project root directory.
- Infrastructure Launch:** Execute the provided shell script.

```
# Option A: Using the unified launcher
./run.sh --docker

# Option B: Using Docker Compose directly
cd infra
docker compose up --build -d
```

- 4. **Database Seeding:** The backend container will automatically apply migrations and seed initial data.
- 5. **Local Frontend Execution (Optional):**

```
./run.sh --search
```

Access Points

Service	URL	Notes
Frontend UI (SMS + Nevai)	http://localhost	React SPA served via Nginx
FastAPI Swagger Docs	http://localhost:8000/docs	Interactive API documentation
Grafana Metrics	http://localhost:3000	User: admin Pass: sms_grafana_2026
Prometheus	http://localhost:9090	Raw metrics explorer

10. Demo Credentials & Roles

The system implements strict **Role-Based Access Control (RBAC)**. Try logging in with these seeded accounts:

Role	Email	Password	Access Level
Admin	admin@sms.edu	Admin@2026	Full Access: User Management, Audit Logs, System Monitoring
Instructor	dr.ahmed@sms.edu	Instructor@2026	Staff Access: View/Manage Students
Student	ahmed.al-hassan0@student.sms.edu	Student@2026	Restricted Access: View/Update Own Profile Only

11. Core API Endpoints

Method	Path	Description	Access
POST	/api/v1/auth/register	Register new user (Strong password enforced)	Public
POST	/api/v1/auth/login	Login and obtain JWT token (set as HttpOnly cookie)	Public
GET	/api/v1/auth/users	List all users with profiles	ADMIN

Method	Path	Description	Access
PUT	/api/v1/auth/users/:id/role	Update a user's role	ADMIN
GET	/api/v1/students	List students (Cached, Paginated, Filterable)	ADMIN/STAFF
GET	/api/v1/students/:id	Get single student detail (Cached)	ADMIN/STAFF
POST	/api/v1/students	Create new student record	ADMIN/STAFF
PUT	/api/v1/students/:id	Update student record (triggers cache invalidation)	ADMIN/STAFF
DELETE	/api/v1/students/:id	Delete student record (triggers cache invalidation)	ADMIN
GET	/api/v1/admin/audit-logs	Paginated system audit trail (JSON Diff)	ADMIN
GET	/api/v1/admin/system-stats	Live analytical data	ADMIN
GET	/api/v1/search/*	Aggregated proxy endpoints for Nevai Search	Authenticated

12. Appendix: Data & Artifact Availability

Artifact	Path
System Launcher	run.sh
Master Development Checklist	todo.txt
Infrastructure Orchestration	infra/docker-compose.yml
Database Schema (Models)	fastapi_backend/app/models/
Telemetry Configuration	infra/prometheus/prometheus.yml
Grafana Dashboard	infra/grafana/dashboard.json
Frontend State Stores	5A_Search/src/store/
Frontend Security Utils	5A_Search/src/utils/security.ts
Application Logs	fastapi_backend/logs/app.log

13. Rubric & Status Table

Grading Criteria	Status	Quality & Notes
JWT Authentication (3)	Complete	Secure refresh rotation. Tokens abstracted to HttpOnly cookies. Strong password enforcement.
Database Design & ORM (2)	Complete	3NF PostgreSQL schema. AsyncPG. Audit trails with JSON diff.
Validation (3)	Complete	Pydantic boundary validation rejects malformed requests (HTTP 422).
Clean Code & Architecture (2)	Complete	Router/Service/Repository separation. DRY patterns. Strict typing.

Grading Criteria	Status	Quality & Notes
Discussion (8)	Complete	Addressed in <i>Results & Discussion</i> , including limitations and future scope.
Logging & Monitoring (3)	Complete	Loguru structured JSON logs. Prometheus/Grafana pipeline. Custom Admin Dashboard.
Redis Caching (2)	Complete	Cache-Aside pattern. Strict invalidation on POST/PUT/DELETE .
API Testing (1)	Complete	84% coverage via Pytest. 37 test cases covering auth, RBAC, CRUD.
Git Workflow (1)	Complete	Clean commit history, logical directory structure.
Frontend (2-Bonus)	Complete	React/Vite SPA. Zustand state. Tailwind layouts. Nevai AI Search integration.
Docker (2-Bonus)	Complete	Multi-container orchestration. Zero-config deployments.

14. Key Innovations / Highlights

- **Innovation 1: Privacy-Preserving Local AI**

By utilizing `web-llm` and WebAssembly, the Nevai search engine runs complex LLMs entirely within the client browser. This enables AI-powered summarization and Q&A over 20+ aggregated search sources without transmitting user data to external servers.

- **Innovation 2: Secure Token Management Strategy**

Abstracts tokens from vulnerable `localStorage` to `HttpOnly` cookies, neutralizing XSS/CSRF attack vectors while maintaining automated refresh rotation.

- **Innovation 3: Predictive Cache Invalidation**

Redis Cache-Aside pattern instantly invalidates related entity lists upon data mutation, ensuring blazingly fast reads with zero stale data.

- **Innovation 4: Type-Safe Data Boundaries**

Strict Pydantic schemas guarantee runtime type safety at the API boundary, eradicating injection vulnerabilities and runtime errors.

- **Innovation 5: Multi-Source Search Aggregation**

Seamless integration of 20+ APIs (Academic, Social, General Knowledge) categorized into distinct search verticals (Web, News, Weather, Definitions) through a unified, minimalist UI.

- **Innovation 6: Advanced Auditing**

Dedicated `audit_logs` table tracking all mutations with JSON diffing for a complete forensic trail.

- **Innovation 7: Zero-Latency State Persistence**

Zustand `localStorage` persistence enables offline-capable productivity tools (Notes/Kanban/Calendar) without backend round-trips.

15. Individual Discussion Prep & Team Instructions

Tech Stack Decision: FastAPI + Python

1. **Pydantic Validation:** Request and response schemas enforced at runtime.
2. **Async Endpoints:** FastAPI supports async I/O for scalable request handling.
3. **SQLAlchemy ORM:** Async sessions ensure clean separation of concerns.

Authentication Strategy

- **JWT via HttpOnly Cookies:** Access/refresh tokens issued with expiry and verified by FastAPI dependency. RBAC ensures students access only their profiles; admins have full CRUD.

Database, Caching & Search Architecture

- **PostgreSQL (3NF):** Normalized `User` → `Student` relationship.
- **Cache-Aside (Redis):** `GET` endpoints cached; updates/deletes invalidate keys automatically.
- **Client-Side AI:** `web-llm` handles search synthesis without backend AI load.

Team Instructions & Whiteboarding

- Review `app/api/auth.py` and `app/db/database.py` to whiteboard the JWT cycle and async DB sessions.
- Review `5A_Search/src/lib/api.ts` to trace how 20+ external APIs are proxied/intercepted and passed to the local Wasm AI.
- **Demo Tip:** Log in as Admin to showcase User Directory/Audit Logs, then perform a Nevai Web Search to demonstrate the local AI summarization feature.

16. System Overview: Architecture & Data Flow

The application utilizes a three-layer decoupled architecture:

1. **Frontend Layer:** React 19 SPA (Vite). Manages UI state (Zustand), local AI inference (`web-llm`), and client-side routing.
2. **Backend Service Layer:** FastAPI (Python). Handles business logic, SMS data validation, secure auth flows, and search API proxying.
3. **Infrastructure & Persistence Layer:** PostgreSQL (ACID store), Redis (Cache), Prometheus/Grafana (Monitoring). Orchestrated via Docker.

Example Data Flow: Nevai AI Search Event

1. User inputs a query into the Nevai UI (`5A_Search/src/components/SearchHome.tsx`).
2. Frontend dispatches parallel requests to the relevant external APIs (or backend proxies) based on selected categories.
3. JSON results are returned to the browser.
4. The `web-llm` module (running locally in Wasm) ingests the JSON context and generates a human-readable summary/answer.
5. The UI renders tabbed results alongside the AI Summary Card (`AISummaryCard.tsx`). No backend AI compute is used.

Output:

```
[SMS] Starting SMS in Docker mode...
[✓] Generated secure JWT secrets
[✓] Initializing infrastructure volumes
[✓] Starting PostgreSQL database and Redis cache services
```

```
[✓] Waiting for database readiness...
[✓] Applying SQLAlchemy Alembic migrations...
[✓] Seeding initial database records...
[✓] Starting FastAPI backend and React frontend...
[✓] All Docker services running! System is online.
```

Step 2: Start Frontend Locally (Dev Mode)

```
./run.sh --search
```

Output:

```
[SMS] Starting 5A_Search application locally...
VITE v6.3.5 ready in 324 ms
→ Local:   http://localhost:5173/
→ Network: use --host to expose
```

Step 3: Halt System

```
./run.sh --stop
```

Output:

```
[SMS] Stopping SMS services...
[✓] Stopping frontend container...
[✓] Stopping backend container...
[✓] Stopping db and redis containers...
[✓] All services safely stopped.
```

21. Frontend Project Overview: 5A_Search

Architecture Deep Dive

The frontend embraces a modern SPA philosophy using React 19 and Vite. It combines the Student Management System interfaces with the Nevai Search Engine into a unified, highly responsive application.

Data Flow:

- User Interaction:** User interacts with UI (e.g., executing a search, updating a student record).
- Global State (Zustand):** Centralized state manages auth flows, local Kanban tasks, and UI preferences. Offline tools use `persist` middleware.
- API Interception Layer (`src/lib/api.ts`):** Configured with `withCredentials: true` for seamless `HttpOnly` JWT attachment. For Nevai, it routes requests to 20+ external sources or backend proxies.
- Local AI Processing:** Search query results are fed into the `web-llm` Wasm instance, generating instant, private AI summaries without backend dependency.
- Dynamic Routing:** React Router enforces UX-level role guards. Expired tokens (`401`) automatically trigger redirect to `LoginPage.tsx`.

User Roles in the UI

- **Admin:** Full UI unlock (Users, Audit Logs, Monitoring, SMS CRUD, Nevai).
- **Instructor:** Pruned navigation (Courses, Students, Nevai).
- **Student:** Restricted view (Personal profile, Nevai).

Security Note: UI pruning is strictly for UX. The backend enforces RBAC as the absolute source of truth, returning `HTTP 403` on unauthorized API calls regardless of client-side state.